

Predicting Promotion Rates through Logistic Regression Modeling

```
In [86]: ► ## import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
import sklearn

## library versions
print('pandas version:', pd.__version__)
print('numpy version:', np.__version__)
print('seaborn version:', sns.__version__)
print('sklearn version:', sklearn.__version__)
```

```
pandas version: 1.0.3
numpy version: 1.18.1
seaborn version: 0.11.2
sklearn version: 0.24.2
```

```
In [87]: ► ## ignore warnings
import warnings
warnings.filterwarnings('ignore')
warnings.simplefilter('ignore')

## print option
pd.set_option('display.max_columns', None)
```

```
In [151]: ► ## Load dataset
df = pd.read_csv("train.csv")
df.head()
```

Out[151]:

	employee_id	department	region	education	gender	recruitment_channel	no_of_trainings	age	previous
0	65438	Sales & Marketing	region_7	Master's & above	f	sourcing	1	35	
1	65141	Operations	region_22	Bachelor's	m	other	1	30	
2	7513	Sales & Marketing	region_19	Bachelor's	m	sourcing	1	34	
3	2542	Sales & Marketing	region_23	Bachelor's	m	other	2	39	
4	48945	Technology	region_26	Bachelor's	m	other	1	45	

```
In [152]: ► ## scrub df
df.replace([np.inf, -np.inf], np.nan, inplace=True)
```

```
In [153]: ► df = df.dropna()
```

EDA

1. Employee ID is PII and not necessary for the model

2. Demographic data includes gender, age, and education
3. Work-related data make up the remainder of the fields

```
In [114]: >>> ## df dimensions
df.shape
```

```
Out[114]: (48660, 13)
```

```
In [115]: >>> ## datatypes
df.dtypes
```

```
Out[115]: employee_id      int64
department    object
region        object
education     object
gender        object
recruitment_channel object
no_of_trainings int64
age           int64
previous_year_rating float64
length_of_service int64
awards_won?    int64
avg_training_score int64
is_promoted    int64
dtype: object
```

```
In [94]: >>> ## summary stats
df.describe()
```

```
Out[94]:
```

	employee_id	no_of_trainings	age	previous_year_rating	length_of_service	awards_won?	avg_training_score
count	48660.000000	48660.000000	48660.000000	48660.000000	48660.000000	48660.000000	48660.000000
mean	39169.271681	1.251993	35.589437	3.337526	6.311157	0.02314	3.337526
std	22630.461554	0.604994	7.534571	1.257922	4.20476	0.15035	1.257922
min	1.000000	1.000000	20.000000	1.000000	1.000000	0.000000	1.000000
25%	19563.500000	1.000000	30.000000	3.000000	3.000000	0.000000	3.000000
50%	39154.000000	1.000000	34.000000	3.000000	5.000000	0.000000	3.000000
75%	58788.250000	1.000000	39.000000	4.000000	8.000000	0.000000	4.000000
max	78298.000000	10.000000	60.000000	5.000000	37.000000	1.000000	5.000000

```
In [95]: >>> ## summary stats - non-numerical
df.describe(include = ['O'])
```

```
Out[95]:
```

	department	region	education	gender	recruitment_channel
count	48660	48660	48660	48660	48660
unique	9	34	3	2	3
top	Sales & Marketing	region_2	Bachelor's	m	other
freq	14239	10811	33404	33852	27017

Numerical Visualizations Observations

Length of Service

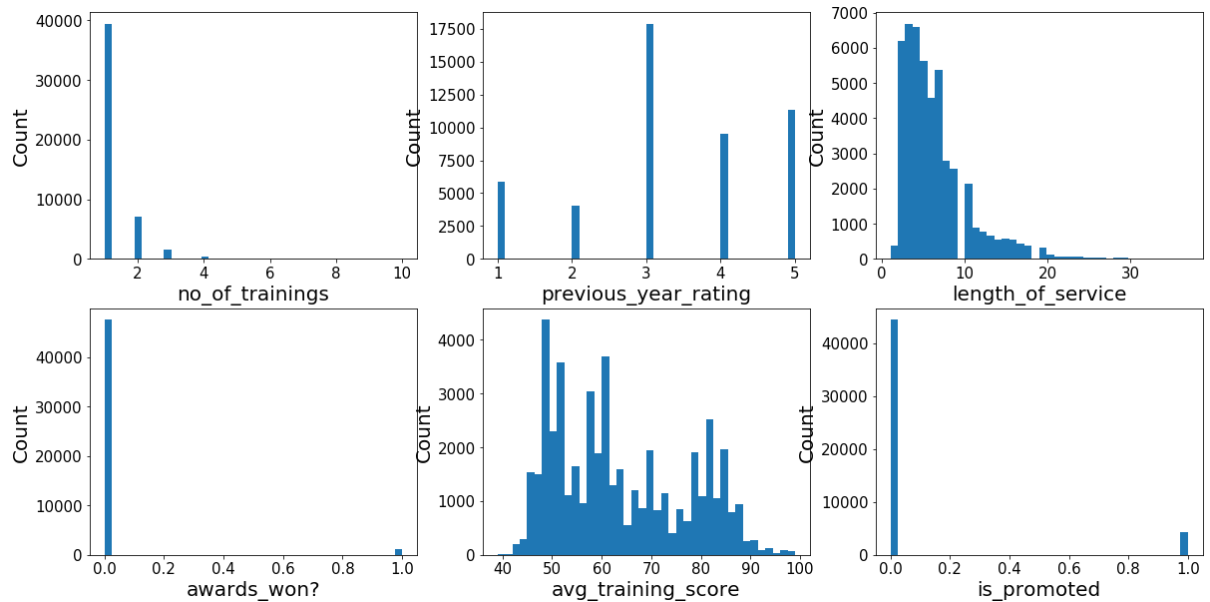
1. Mean(5.8) > median(5.0)
2. Length of service ranges from 1 yr to 37 yrs
3. The length of service of most employees is centered between 1 and 6 yrs

Average training score

1. Mean(63.38) > median(60.00)
2. Average training score ranges from 39 and 99

```
In [13]: ## numerical data histograms


plt.rcParams['figure.figsize'] = (20, 10)
fig, axes = plt.subplots(nrows = 2, ncols = 3)
num_features = ['no_of_trainings', 'previous_year_rating', 'length_of_service', 'award:
xaxes = num_features
yaxes = ['Count', 'Count', 'Count', 'Count', 'Count', 'Count', 'Count']
axes = axes.ravel()
for idx, ax in enumerate(axes):
    ax.hist(df[num_features[idx]].dropna(), bins=40)
    ax.set_xlabel(xaxes[idx], fontsize=20)
    ax.set_ylabel(yaxes[idx], fontsize=20)
    ax.tick_params(axis='both', labelsize=15)
plt.show()
```




Outliers Observations

1. Average Training Score does not have a noticeable outlier
2. Length of Service returned a slight outlier. Potential max range would be 13

```
In [24]: ► ## outlier check
fig = px.box(df,y='avg_training_score',points='outliers', title='Outliers in Average T
fig.update_layout(hovermode='x')
```



```
In [25]: fig = px.box(df, y='length_of_service', points='outliers', title='Outliers in Length of  
fig.update_layout(hovermode='x')
```



Promoted Visualizations Observations

Department

1. R&D is the smallest department

Education


1. More than 35000 employees hold a bachelor's degree
2. At least 15000 employees have a Master's and Phd

Gender

1. Male employees account for more than 35000 employees in the company
2. The number of female employees is slightly above 15000

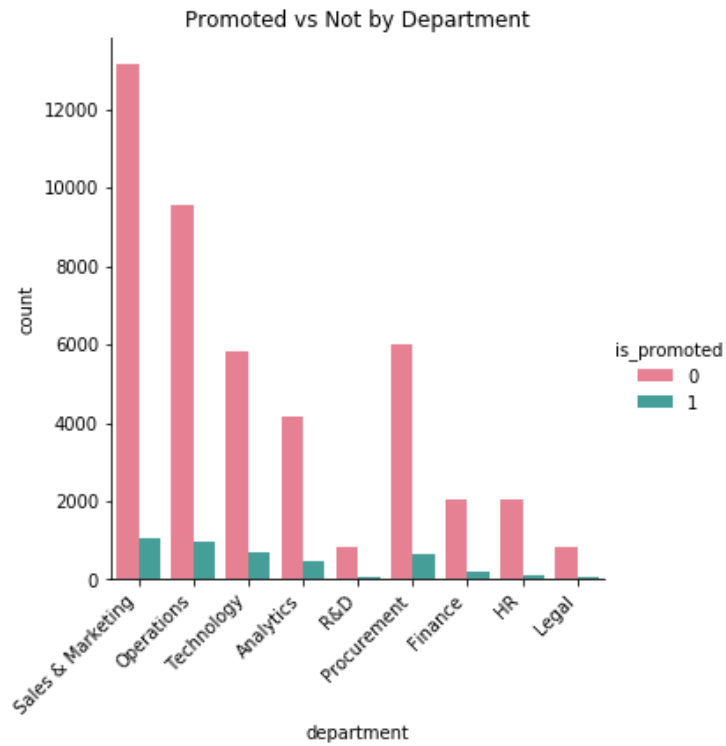
```
In [28]: ► ## department distribution
plt.figure(figsize = (12, 10))
sns.catplot(x = 'department', kind='count', data = df, palette = 'husl')
plt.xticks(rotation = 45, horizontalalignment = 'right')
plt.title('Employees per Department')
plt.show()
```

<Figure size 864x720 with 0 Axes>



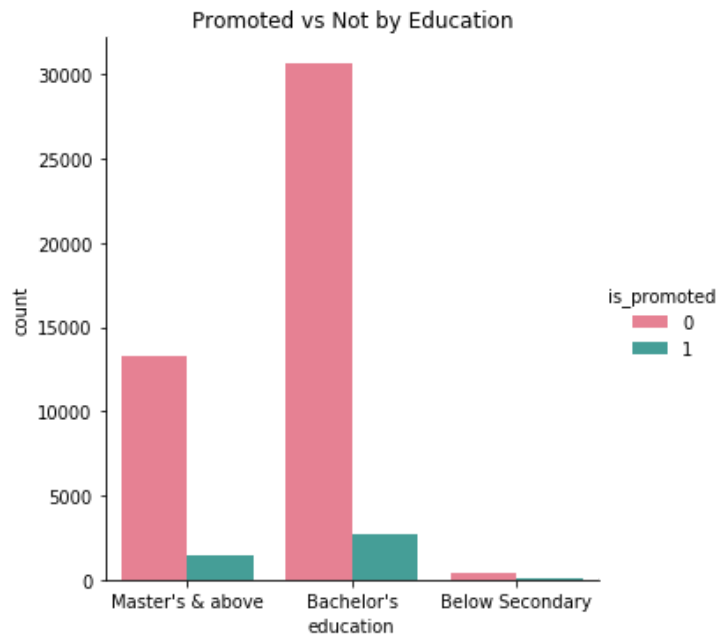
```
In [29]: ▶ ## promoted vs not by department
plt.figure(figsize = (12, 10))
sns.catplot(x = 'department', hue = 'is_promoted', kind = 'count', data = df, palette=
plt.xticks(rotation=45, horizontalalignment='right')
plt.title('Promoted vs Not by Department')
plt.show()
```

<Figure size 864x720 with 0 Axes>



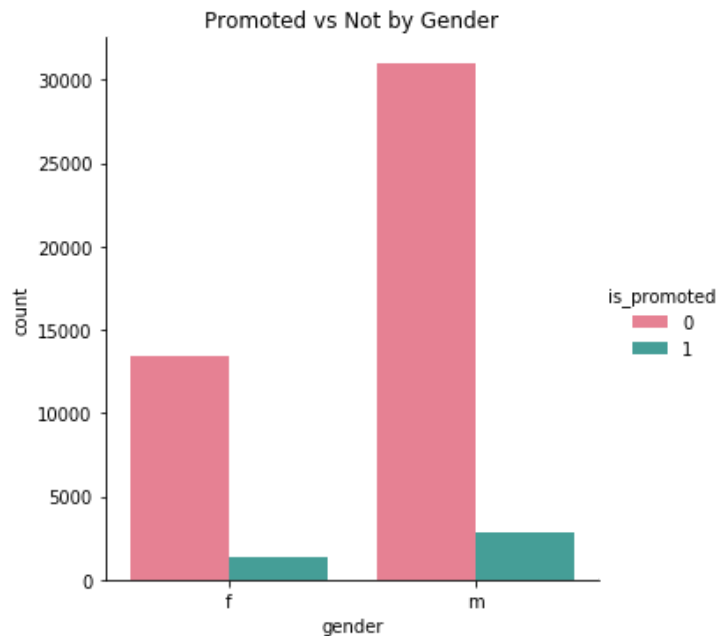
```
In [31]: ► ## promoted vs not by education
plt.figure(figsize=(12, 10))
sns.catplot(x = 'education', hue = 'is_promoted', kind = 'count', data = df, palette = 'h
plt.title('Promoted vs Not by Education')
plt.show()
```

<Figure size 864x720 with 0 Axes>



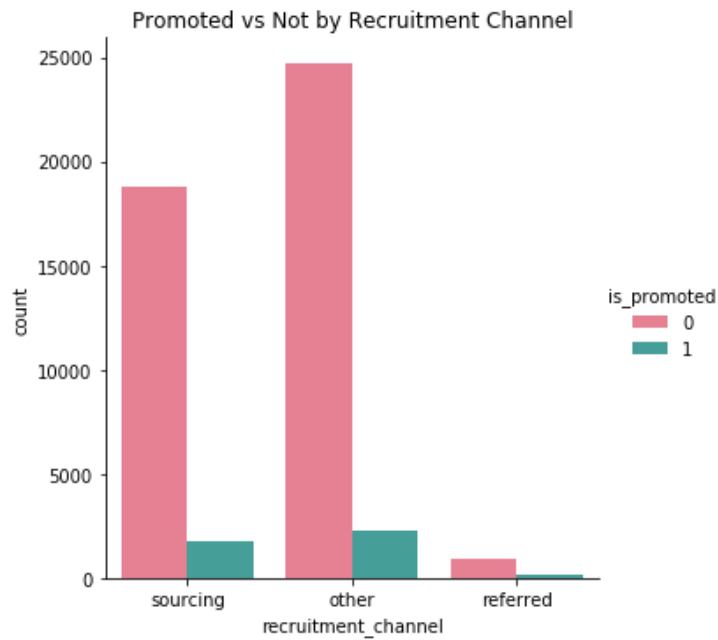
```
In [32]: ► ## promoted vs not by gender
plt.figure(figsize=(12, 10))
sns.catplot(x = 'gender', hue = 'is_promoted', kind = 'count', data = df, palette = 'h
plt.title('Promoted vs Not by Gender')
plt.show()
```

<Figure size 864x720 with 0 Axes>



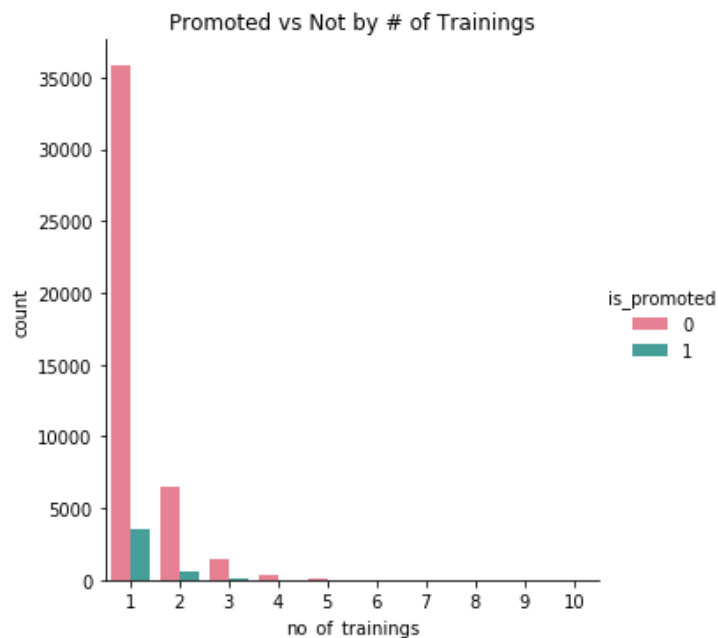

```
In [33]: ▶ ## promoted vs not by recruitment channel
plt.figure(figsize=(12, 10))
sns.catplot(x = 'recruitment_channel', hue = 'is_promoted', kind = 'count', data = df,
plt.title('Promoted vs Not by Recruitment Channel')
plt.show()
```

<Figure size 864x720 with 0 Axes>



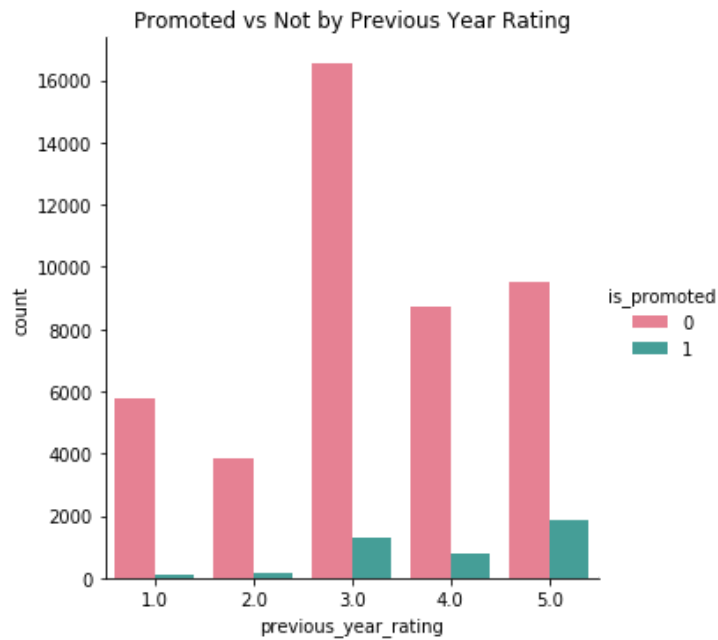
```
In [36]: ▶ ## promoted vs not by # of trainings
plt.figure(figsize=(12, 10))
sns.catplot(x = 'no_of_trainings', hue = 'is_promoted', kind = 'count', data = df, palette='magma')
plt.title('Promoted vs Not by # of Trainings')
plt.show()
```

<Figure size 864x720 with 0 Axes>




```
In [35]: ▶ ## promoted vs not by previous year rating
plt.figure(figsize=(12, 10))
sns.catplot(x = 'previous_year_rating', hue = 'is_promoted', kind = 'count', data = df)
plt.title('Promoted vs Not by Previous Year Rating')
plt.show()
```

<Figure size 864x720 with 0 Axes>

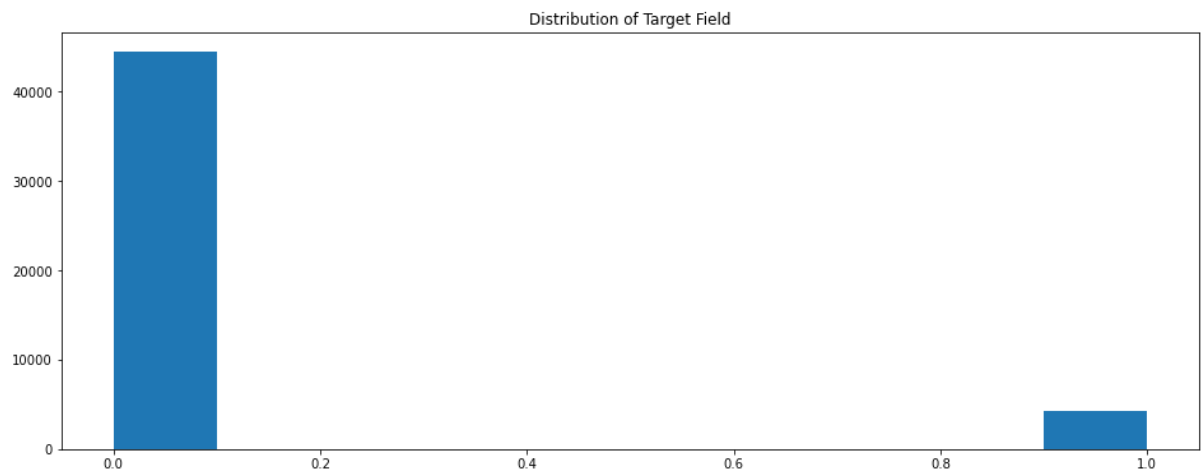


```
In [33]: ▶ ## target field vs workforce data

fig = plt.figure(figsize=(10,5) )
fig.add_subplot(1,3,1)
ar_6 = sns.boxplot(x=df["is_promoted"],y=df["length_of_service"])
fig.add_subplot(1,3,2)
ar_6 = sns.boxplot(x=df["is_promoted"],y=df["avg_training_score"])
fig.add_subplot(1,3,3)
ar_6 = sns.boxplot(x=df["is_promoted"],y=df["previous_year_rating"])
plt.tight_layout()
plt.show()
```




```
In [34]: ► ## distribution of target field
fig = plt.figure(figsize =(16, 6))
plt.hist(df['is_promoted'])
plt.title('Distribution of Target Field', loc = 'center', fontsize = 12)
plt.show()
```



```
In [116]: px.imshow(df.corr(), text_auto= True, title='Correlation Between the Variables in the Model')
```

Correlation Between the Variables in the Model



```
In [156]: ## create dummy variables

df["gender"] = df["gender"].apply(lambda x: 1 if x=="m" else 0)

cols = df.select_dtypes(["object"]).columns
ds = pd.get_dummies(df[cols],drop_first=True)
ds

df = pd.concat([df,ds],axis=1)

## drop original columns
df.drop(cols,axis=1,inplace=True)
```

```
In [155]: ## drop employee_id

df = df.drop('employee_id', axis=1)
df.head()
```

```
Out[155]:
```

	department	education	gender	recruitment_channel	no_of_trainings	age	previous_year_rating	length_of_s
0	Sales & Marketing	Master's & above	f	sourcing	1	35	5.0	
1	Operations	Bachelor's	m	other	1	30	5.0	
2	Sales & Marketing	Bachelor's	m	sourcing	1	34	3.0	
3	Sales & Marketing	Bachelor's	m	other	2	39	1.0	
4	Technology	Bachelor's	m	other	1	45	3.0	

```
In [37]: df.shape
```

```
Out[37]: (48660, 53)
```

```
In [38]: df.head()
```

```
Out[38]:
```

	gender	no_of_trainings	age	previous_year_rating	length_of_service	awards_won?	avg_training_score	is_
0	0	1	35	5.0	8	0	49	
1	1	1	30	5.0	4	0	60	
2	1	1	34	3.0	7	0	50	
3	1	2	39	1.0	10	0	50	
4	1	1	45	3.0	2	0	73	

Logistic Regression Model

```
In [140]: > ## import library
from sklearn.model_selection import train_test_split

## split data
y = df.pop("is_promoted")
X = df

X_train,X_test,y_train,y_test = train_test_split(X,y,random_state=42,train_size=0.8)

print("train size X : ",X_train.shape)
print("train size y : ",y_train.shape)
print("test size X : ",X_test.shape)
print("test size y : ",y_test.shape)

train size X : (38928, 52)
train size y : (38928,)
test size X : (9732, 52)
test size y : (9732,)
```

```
In [142]: > from sklearn.preprocessing import StandardScaler

scale = StandardScaler()

X_train = scale.fit_transform(X_train)
X_test = scale.transform(X_test)
```

```
In [143]: > ## distribution values check
y_train.value_counts(normalize=True)
```

```
Out[143]: 0    0.913558
          1    0.086442
          Name: is_promoted, dtype: float64
```

```
In [144]: > ## import library
from sklearn.linear_model import LogisticRegression

## add class weight to address 90/10 promoted split
lr = LogisticRegression(class_weight={0:0.1,1:0.9})
lr.fit(X_train,y_train)
```

```
Out[144]: LogisticRegression(class_weight={0: 0.1, 1: 0.9})
```

```
In [145]: > ## set base model
base_model = lr
y_pred_base_model = base_model.predict(X_test)
pred_prob = base_model.predict_proba(X_test)
```

```
In [146]: > ## iomport library
from sklearn.metrics import confusion_matrix


cm = confusion_matrix(y_test, y_pred_base_model)
```

```
In [147]: ## confusion matrix of base model

df1 = pd.DataFrame(columns=["0","1"], index= ["0","1"], data= cm )

f,ax = plt.subplots(figsize=(6,6))

sns.heatmap(df1, annot=True,cmap="Greens", fmt= '.0f',
            ax=ax,linewidths = 5, cbar = False,annot_kws={"size": 12})
plt.xlabel("Predicted Label")
plt.xticks(size = 12)
plt.yticks(size = 12, rotation = 0)
plt.ylabel("True Label")
plt.title("Confusion Matrix Base Model", size = 12)
plt.show()
```



```
In [62]: > ## import library
from sklearn.metrics import roc_curve


## roc curve for classes
fpr = {}
tpr = {}
thresh = {}

n_class = 2

for i in range(n_class):
    fpr[i], tpr[i], thresh[i] = roc_curve(y_test, pred_prob[:,i], pos_label=i)

plt.plot(fpr[0], tpr[0], linestyle='--',color='orange', label='Class 0 ')
plt.plot(fpr[1], tpr[1], linestyle='--',color='green', label='Class 1 ')

plt.title('Multiclass ROC curve')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive rate')
plt.legend(loc='best')
plt.savefig('Multiclass ROC');
```



```
In [148]: > from sklearn.metrics import accuracy_score
accuracy_score(y_test, y_pred_base_model)
```

Out[148]: 0.7950061652281134

Feature Engineering

```
In [149]: > ## combine 'awards_won' and 'previous_year_rating'
## combine 'avg_training_score' and 'no_of_trainings'

#Creating a sum metric column
df['sum_metric'] = df['awards_won?'] + df['previous_year_rating']

# creating a total score column
df['total_score'] = df['avg_training_score'] * df['no_of_trainings']
```



```
In [154]: ▶ ## drop region  
df = df.drop(['region'], axis = 1)
```

```
In [157]: ▶ df.columns
```

```
Out[157]: Index(['gender', 'no_of_trainings', 'age', 'previous_year_rating',  
                'length_of_service', 'awards_won?', 'avg_training_score', 'is_promoted',  
                'department_Finance', 'department_HR', 'department_Legal',  
                'department_Operations', 'department_Procurement', 'department_R&D',  
                'department_Sales & Marketing', 'department_Technology',  
                'education_Below Secondary', 'education_Master's & above',  
                'recruitment_channel_referred', 'recruitment_channel_sourcing'],  
              dtype='object')
```

Label Encode

```
In [158]: ▶ df.select_dtypes('object').head()
```

```
Out[158]:  
_____  
0  
1  
2  
3  
4
```

```
In [162]: ▶ df.head()
```

```
Out[162]:
```

	gender	no_of_trainings	age	previous_year_rating	length_of_service	awards_won?	avg_training_score	is_
0	0	1	35	5.0	8	0	49	
1	1	1	30	5.0	4	0	60	
2	1	1	34	3.0	7	0	50	
3	1	2	39	1.0	10	0	50	
4	1	1	45	3.0	2	0	73	

2nd Logistic Regression Model

```
In [163]: ▶ #split target from features  
  
y = df['is_promoted']  
x = df.drop(['is_promoted'],axis=1)
```

```
In [165]: > ## split data
X_train,X_test,y_train,y_test = train_test_split(X,y,random_state=42,train_size=0.8)

print("train size X : ",X_train.shape)
print("train size y : ",y_train.shape)
print("test size X : ",X_test.shape)
print("test size y : ",y_test.shape)
```

```
train size X : (38928, 54)
train size y : (38928,)
test size X : (9732, 54)
test size y : (9732,)
```

```
In [166]: > scale = StandardScaler()

X_train = scale.fit_transform(X_train)
X_test = scale.transform(X_test)
```

```
In [167]: > ## distribution values check
y_train.value_counts(normalize=True)
```

```
Out[167]: 0    0.913558
          1    0.086442
          Name: is_promoted, dtype: float64
```

```
In [168]: > ## add class weight to address 90/10 promoted split
lr = LogisticRegression(class_weight={0:0.1,1:0.9})
lr.fit(X_train,y_train)
```

```
Out[168]: LogisticRegression(class_weight={0: 0.1, 1: 0.9})
```

```
In [169]: > ## set base model
base_model = lr
y_pred_base_model = base_model.predict(X_test)
pred_prob = base_model.predict_proba(X_test)
```

```
In [170]: > cm = confusion_matrix(y_test, y_pred_base_model)
```

```


In [172]: ► ## confusion matrix of base model

df1 = pd.DataFrame(columns=["0","1"], index= ["0","1"], data= cm )

f,ax = plt.subplots(figsize=(6,6))

sns.heatmap(df1, annot=True,cmap="Greens", fmt= '.0f',
            ax=ax,linewidths = 5, cbar = False,annot_kws={"size": 12})
plt.xlabel("Predicted Label")
plt.xticks(size = 12)
plt.yticks(size = 12, rotation = 0)
plt.ylabel("True Label")
plt.title("Confusion Matrix Base Model", size = 12)
plt.show()

```



```

In [173]: ► accuracy_score(y_test, y_pred_base_model)

```

Out[173]: 0.7949034114262228

Feature Evaluation

```
In [186]: from sklearn.datasets import make_regression
from sklearn.linear_model import LinearRegression
from matplotlib import pyplot

X, y = make_regression(n_samples=1000, n_features=20, n_informative=5, random_state=1)

lr = LinearRegression()

lr.fit(X, y)

importance = lr.coef_

for i,v in enumerate(importance):
    print('Feature: %0d, Score: %.5f' % (i,v))
pyplot.bar([x for x in range(len(importance))], importance)
pyplot.title('Feature Importance Scores')
pyplot.show()
```

```
Feature: 0, Score: 0.00000
Feature: 1, Score: -0.00000
Feature: 2, Score: -0.00000
Feature: 3, Score: 0.00000
Feature: 4, Score: -0.00000
Feature: 5, Score: 0.00000
Feature: 6, Score: 9.61372
Feature: 7, Score: 0.00000
Feature: 8, Score: 30.51944
Feature: 9, Score: 0.00000
Feature: 10, Score: 0.00000
Feature: 11, Score: 0.00000
Feature: 12, Score: 48.38204
Feature: 13, Score: 0.00000
Feature: 14, Score: -0.00000
Feature: 15, Score: -0.00000
Feature: 16, Score: 39.99774
Feature: 17, Score: -0.00000
Feature: 18, Score: -0.00000
Feature: 19, Score: 70.86224
```

